

Leihan Wang

 leihanwa@buffalo.edu |  <https://github.com/LWang1101> |  [Leihan Wang - Google Scholar](#)

EDUCATION

School of Engineering and Applied Sciences

University at Buffalo, SUNY (Buffalo, New York, The United States)

Bachelor of Science in **Computer Science**

August 2024 - Present

College of Artificial Intelligence and Automation

Hohai University (Nanjing, Jiangsu, China)

Bachelor of Engineer in **Artificial Intelligence**

August 2022 - July 2024

I am an undergraduate student from *Hohai University* to *University at Buffalo* 😊

RESEARCH EXPERIENCE

University at Buffalo, SUNY (Buffalo, New York, The United States)

August 2026 – April 2026

Research Student, hosted by Prof. [Yaxiong Xie](#)

- Conducted systems research on **mobile SoC architectures for on-device LLM inference**, focusing on heterogeneous CPU/GPU/NPU execution, Qualcomm Hexagon/HVX vector processing, memory hierarchy, and data movement across mobile accelerators.
- Investigated modern **NPU-based LLM inference systems including NPUGen and ShadowNPU**, analyzing architectural techniques such as KV-cache recomputation, sparse attention, accelerator offloading, and the limitations imposed by NPU programmability and operator support.
- Explored **architecture-aware optimization opportunities for mobile LLM inference** by connecting prefill, decoding, and attention workloads with hardware characteristics including SIMD parallelism, cache/VTM behavior, memory bandwidth, instruction scheduling, and power/thermal constraints.

Clemson University (Clemson, South Carolina, The United States)

May 2026 - September 2026

Remote Research Student, hosted by Prof. [Xiaoyong \(Brian\) Yuan](#)

- Developed a multimodal foundation-model alignment framework for drug discovery, integrating molecular, protein, textual, and biomedical representations.
- Built fusion strategies to align heterogeneous foundation models and improve molecule–target association, drug property prediction, and biomedical reasoning.
- Incorporated domain-specific chemical and biological priors to enhance cross-modal representation learning and foundation-model-based drug discovery.

The University of Georgia (Athens, Georgia, The United States)

December 2025 – February 2026

Remote Research Student, hosted by Prof. [Zhenlin An](#)

- Developing a PE-AV model and beamforming-based audio-visual perception system for AR glasses.
- Building an AR audio-visualization framework that localizes sound sources and visually maps audio signals to corresponding sound-emitting objects in the environment.
- Incorporating visual priors to improving sound source association, spatial audio understanding, and object-level audio-visual reasoning.

- Research Student, hosted by Prof. [Haipeng Cai](#)
- Developed an enhanced BODE-GEN framework for LLM-driven program synthesis, using Bayesian Optimization to discover task-adaptive prompts and exemplars from a subset of HumanEval+ problems.
 - Integrated an auxiliary LLM with a Bayesian-optimized prompt memory to select task-specific prompt–exemplar combinations for unseen coding problems, improving generalization beyond a single globally optimized prompt.
 - Reduced reliance on expensive per-task Bayesian Optimization by optimizing only 20 HumanEval+ tasks offline and transferring the learned prompt memory to the remaining 144 tasks for single-generation pass@1 evaluation.

- Research assistant, hosted by Prof. [Zaipeng Xie](#)
- Developed **FedSTAR** and **FedHAN**, federated learning frameworks addressing **spatiotemporal/data heterogeneity, model heterogeneity, and label noise** across decentralized trajectory prediction and heterogeneous client models. FedSTAR specifically combines alignment, drift control, and performance-aware aggregation. FedHAN focuses on reliability under heterogeneous architectures and noisy labels.
 - Designed robust local and global optimization mechanisms, including **two-stage contrastive alignment, adaptive model-drift regularization, reliability-weighted aggregation, consistency-based sample weighting, and prototype-based classifier refinement**, improving stability and knowledge sharing across heterogeneous clients.
 - Evaluated the frameworks across **ETH/UCY and MNIST/CIFAR-10/CIFAR-100**, with FedSTAR achieving the lowest average ADE/FDE among federated baselines under heterogeneous trajectory partitions and maintaining performance under dynamic client participation, while FedHAN improved accuracy by up to **5.76% on CIFAR-10 and 4.67% on CIFAR-100** under heavy label noise.

PROJECT EXPERIENCE

- Project Engineer
- Developed **Siemens S7-300 PLC control logic using STEP 7** to coordinate multiple water pumps based on real-time flow, pressure, and process parameters, supporting automated filtration and reliable water treatment operations.
 - Implemented **automated chemical dosing, pump rotation, fault detection, and backup-pump switching**, enabling dynamic process control and stable unattended operation under equipment or sensor failures.
 - Applied **control engineering, industrial automation, and chemical process knowledge** to support system design, technical documentation, and preparation of a municipal water-treatment bidding project.

PUBLICATIONS

[2] (**IEEE BigData 2025**) Zaipeng Xie, Zishu Zhou, Xuanyao Jie, Xing Gao, Han Xu, and **Leihan Wang**, “FedHAN: Robust Federated Learning under Model Heterogeneity and Label Noise.” Accepted to *2025 IEEE International Conference on Big Data* (IEEE BigData 2025), 2025.

[1] (**EAI CollaborateCom 2025**) Zaipeng Xie, Dingxu Sun, Peixin Li, Ziyang Ye, and **Leihan Wang**, “FedSTAR: A Federated Learning Framework for Reliable Trajectory Prediction under Spatiotemporal Heterogeneity.” Accepted to *21st EAI International Conference on Collaborative Computing: Networking, Applications and Worksharing* (EAI CollaborateCom), 2025.

SELECTED COURCES GRADES

• Algorithms and Complexity	98/100	Artificial Intelligence	91/100
• Algorithms and Data Structures	96/100	Introduction to UAV	96/100